



RAJALAKSHMI
ENGINEERING COLLEGE
An AUTONOMOUS Institution
Affiliated to ANNA UNIVERSITY, Chennai

FUNDAMENTALS OF DATA SCIENCE

Laboratory Record Notebook

Name: AWINTHIKA SANTHANAM

Year / Branch / Section: IInd Year - CSE 'A' - B.E

University Register No: 230701048

College Roll No: 230701048

Semester: III

Academic Year: 2024-2025

```
In [ ]: # Name of the Experiment : Pandas Built in function
        # EX NO : 01
        # Register Number : 230701048
        # Name : AWINTHIKA SANTHANAM
```

```
In [1]: import pandas as pd
import numpy as np
# --- Part 1: Pandas --- #

# Create a sample DataFrame with multiple columns and rows
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eva'],
    'Age': [25, 30, 35, 40, 45],
    'Salary': [50000, 60000, 70000, 80000, 90000],
    'Department': ['HR', 'IT', 'Finance', 'Marketing', 'Sales']
}

df = pd.DataFrame(data) # Display the DataFrame
print("Original DataFrame:")
print(df)
print("-" * 50)
# Access specific rows and columns using `iloc` (index-based slicing)
print("Sliced DataFrame using iloc (index-based slicing):")
print(df.iloc[1:4]) # Slicing rows 1 to 3 (0-indexed)
print("-" * 50)
# Access specific rows and columns using `loc` (label-based slicing)
print("Sliced DataFrame using loc (label-based slicing):")
print(df.loc[1:3]) # Slicing rows 1 to 3 (inclusive)
print("-" * 50)
# Access specific column(s)
print("Accessing 'Age' column:")
print(df['Age'])
print("-" * 50)
# Select multiple columns
print("Accessing multiple columns 'Name' and 'Salary':")
print(df[['Name', 'Salary']])
print("-" * 50)
# Filter DataFrame based on condition
print("Filtering DataFrame where Age > 30:")
print(df[df['Age'] > 30])
print("-" * 50)
# Add a new column with calculated values
df['Salary_in_K'] = df['Salary'] / 1000
print("DataFrame with new 'Salary_in_K' column:")
print(df)
print("-" * 50)

# Number of dimensions (ndim) of the DataFrame
print("Number of dimensions of DataFrame:")
print(df.ndim) # Should return 2 (since it's a DataFrame)
print("-" * 50)
# --- Part 2: NumPy --- # Create a NumPy 2D array
arr = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
# Display the original NumPy array
print("Original NumPy array:")
print(arr)
print("-" * 50)
# Reshaping the array to 1D using reshape()
reshaped_arr = arr.reshape(-1)
print("Reshaped NumPy array (1D):")
print(reshaped_arr)
print("-" * 50)
# Flatten the array using ravel()
raveled_arr = arr.ravel()
print("Raveled NumPy array:")
```

```

print(raveled_arr) print("-" * 50)
# Number of dimensions (ndim) of the NumPy array
print("Number of dimensions of NumPy array:")
print(arr.ndim) # Should return 2 (since it's a 2D array)
print("-" * 50)
# Array slicing (slicing rows and columns)
print("Array slicing - Select rows 1 and 2, columns 0 and 1:")
print(arr[1:3, 0:2]) # Slicing rows 1 to 2, and columns 0 to 1
print("-" * 50)
# Transpose the array
transposed_arr = arr.T
print("Transposed NumPy array:")
print(transposed_arr)
print("-" * 50)

# Operations on NumPy array (addition, multiplication)
print("Element-wise addition (arr + 10):")
print(arr + 10)
print("-" * 50)
print("Element-wise multiplication (arr * 2):")
print(arr * 2)
print("-" * 50)

# --- Part 3: Combining Pandas and NumPy --- #

# Convert DataFrame column to NumPy array
numpy_salary = df['Salary'].to_numpy()
print("Converted 'Salary' column from DataFrame to NumPy array:")
print(numpy_salary)
print("-" * 50)

# Perform a NumPy operation on the 'Salary' column of the DataFrame
new_salaries = numpy_salary * 1.1 # Increase salary by 10%
df['Updated_Salary'] = new_salaries
print("DataFrame with updated salaries (10% increase):")
print(df)
print("-" * 50)

# Create a NumPy array from multiple columns of the DataFrame
salary_dept_arr = df[['Salary', 'Department']].to_numpy()
print("NumPy array from 'Salary' and 'Department' columns of DataFrame:")
print(salary_dept_arr)
print("-" * 50)

# --- End of Program ---

```

OUTPUT

```

Original DataFrame:
   Name  Age  Salary Department
0  Alice   25   50000         HR
1   Bob    30   60000         IT
2  Charlie   35   70000    Finance
3   David   40   80000    Marketing
4    Eva    45   90000     Sales
-----
Sliced DataFrame using iloc (index-based slicing):
   Name  Age  Salary Department
1   Bob    30   60000         IT
2  Charlie   35   70000    Finance
3   David   40   80000    Marketing
-----
Sliced DataFrame using loc (label-based slicing):
   Name  Age  Salary Department
1   Bob    30   60000         IT
2  Charlie   35   70000    Finance
3   David   40   80000    Marketing
-----
Accessing 'Age' column:
0    25
1    30
2    35
3    40
4    45
Name: Age, dtype: int64
-----
Accessing multiple columns 'Name' and 'Salary':
   Name  Salary
0  Alice   50000
1   Bob   60000
2  Charlie  70000

```

```

3 David 80000 4 Eva 90000 -----
----- Filtering DataFrame where Age > 30:
      Name Age Salary Department 2 Charlie 35 70000 Finance
3 David 40 80000 Marketing 4 Eva 45 90000 Sales -----
-----
DataFrame with new 'Salary_in_K' column:
      Name Age Salary Department Salary_in_K
0 Alice 25 50000 HR 50.0
1 Bob 30 60000 IT 60.0
2 Charlie 35 70000 Finance 70.0
3 David 40 80000 Marketing 80.0
4 Eva 45 90000 Sales 90.0 -----
----- Number of dimensions of DataFrame: 2 -----
----- Original NumPy array: [[1 2 3] [4 5 6] [7 8 9]] -----
----- Reshaped NumPy array (1D):
[1 2 3 4 5 6 7 8 9] -----
----- Raveled NumPy array: [1 2 3 4 5 6 7 8 9] -----
----- Number of dimensions of NumPy array: 2 -----
----- Array slicing -----
Select rows 1 and 2, columns 0 and 1: [[4 5] [7 8]] -----
----- Transposed NumPy array: [[1 4 7] [2 5 8] [3 6 9]] ---
----- Element-wise addition (arr + 10):
[[11 12 13] [14 15 16] [17 18 19]] -----
----- Element-wise multiplication (arr * 2): [[ 2 4 6] [ 8 10 12] [14 16 18]] -----
-----
Converted 'Salary' column from DataFrame to NumPy array:
[50000 60000 70000 80000 90000] -----
----- DataFrame with updated salaries (10% increase):
      Name Age Salary Department Salary_in_K Updated_Salary
0 Alice 25 50000 HR 50.0 55000.0
1 Bob 30 60000 IT 60.0 66000.0
2 Charlie 35 70000 Finance 70.0 77000.0
3 David 40 80000 Marketing 80.0 88000.0
4 Eva 45 90000 Sales 90.0 99000.0 -----
-----
NumPy array from 'Salary' and 'Department' columns of DataFrame:
[[50000 'HR'] [60000 'IT'] [70000 'Finance'] [80000 'Marketing']
[90000 'Sales']] -----

```

```

In [3]: import pandas as pd

# Create a DataFrame
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David'],
    'Age': [24, 27, 22, 32],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']}
df = pd.DataFrame(data)

# Display first few rows
print("Head of DataFrame:")
print(df.head())

# Display last few rows
print("\nTail of DataFrame:")
print(df.tail())

# Summary statistics
print("\nSummary Statistics:")
print(df.describe())

# Information about DataFrame
print("\nDataFrame Info:")
df.info()

```

OUTPUT

```
Head of DataFrame:
  Name  Age   City
0  Alice  24  New York
1   Bob  27 Los Angeles
2 Charlie  22   Chicago
3  David  32   Houston
Tail of DataFrame:
  Name  Age   City
0  Alice  24  New York
1   Bob  27 Los Angeles
2 Charlie  22   Chicago
3  David  32   Houston
Summary Statistics:
      Age
count  4.000000
mean   26.250000
std     4.349329
min    22.000000
25%    23.500000
50%    25.500000
75%    28.250000
max    32.000000
DataFrame Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0  Name    4 non-null      object
1  Age     4 non-null      int64
2  City    4 non-null      object
dtypes: int64(1), object(2)
memory usage: 228.0+ bytes
```

```
In [5]: import numpy as np
```

```
# Create a NumPy array
arr = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9])
# Slice array (from index 2 to 7, with a step of 2)
sliced_arr = arr[2:8:2]
print("Sliced Array:", sliced_arr)
```

Sliced Array: [3 5 7]

```
In [7]: # Create a 2D array
arr_2d = np.array([[1, 2, 3], [4, 5, 6]])
# Flatten the array
flat_arr = arr_2d.ravel()
print("Flattened Array:", flat_arr)
```

Flattened Array: [1 2 3 4 5 6]

```
In [9]: # Reshape 1D array into a 3x3 matrix
reshaped_arr = arr.reshape(3, 3)
print("Reshaped Array (3x3):\n", reshaped_arr)
```

Reshaped Array (3x3):
[[1 2 3]
 [4 5 6]
 [7 8 9]]

```
In [11]: # Check the number of dimensions
print("Number of Dimensions:", arr.ndim)
print("Number of Dimensions (2D array):", arr_2d.ndim)
```

Number of Dimensions: 1
Number of Dimensions (2D array): 2

```
In [13]: import numpy as np

# Create an array
arr = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9])
# Array slicing
sliced_arr = arr[2:8:2]
print("Sliced Array:", sliced_arr)

# Ravel (flatten the array)
arr_2d = np.array([[1, 2, 3], [4, 5, 6]])
flat_arr = arr_2d.ravel()
print("Flattened Array:", flat_arr)
# Reshape
reshaped_arr = arr.reshape(3, 3)
print("Reshaped Array (3x3):\n", reshaped_arr)

# Number of dimensions
print("Number of Dimensions (original array):", arr.ndim)
print("Number of Dimensions (2D array):", arr_2d.ndim)
```

```
Sliced Array: [3 5 7]
Flattened Array: [1 2 3 4 5 6]
Reshaped Array (3x3):
[[1 2 3]
 [4 5 6]
 [7 8 9]]
Number of Dimensions (original array): 1
Number of Dimensions (2D array): 2
```

```
In [15]: # Name of the Experiment : Outlier detection
# EX NO : 02
# Register Number : 230701048
#Name : AWINTHIKA SANTHANAM
```

```
In [17]: #sample calculation for low range(lr) , upper range (ur),percentile
import numpy as np
array=np.random.randint(1,100,16) # randomly generate 16 numbers between 1 to 100
array
```

```
Out[17]: array([76, 47, 16, 24, 21, 95, 8, 19, 25, 8, 65, 21, 85, 64, 45, 19])
```

```
In [21]: array.mean()
```

```
Out[21]: 39.875
```

```
In [23]: np.percentile(array,25)
```

```
Out[23]: 19.0
```

```
In [25]: np.percentile(array,50)
```

```
Out[25]: 24.5
```

```
In [27]: np.percentile(array,75)
```

```
Out[27]: 64.25
```

```
In [29]: np.percentile(array,100)
```

```
Out[29]: 95.0
```

```
In [33]: def outDetection(array):
sorted(array)
Q1,Q3=np.percentile(array,[25,75])
IQR=Q3-Q1
lr=Q1-(1.5*IQR)
ur=Q3+(1.5*IQR)
return lr,ur
```


```
In [35]: lr,ur=outDetection(array)
```

```
In [37]: lr,ur
```

```
Out[37]: (-48.875, 132.125)
```


```
In [39]: import seaborn as sns
         %matplotlib inline
         sns.displot(array)
```

Out[39]: <seaborn.axisgrid.FacetGrid at 0x2737269b470>



```
In [41]: sns.distplot(array)
```

Out[41]: <Axes: ylabel='Density'>



```
In [43]: new_array=array[(array>lr) & (array<ur)]
         new_array
```

Out[43]: array([76, 47, 16, 24, 21, 95, 8, 19, 25, 8, 65, 21, 85, 64, 45, 19])

```
In [45]: lr1,ur1=outDetection(new_array)
         lr1,ur1
```


Out[45]: (-48.875, 132.125)

```
In [47]: final_array=new_array[(new_array>lr1) & (new_array<ur1)]
final_array
```

Out[47]: array([76, 47, 16, 24, 21, 95, 8, 19, 25, 8, 65, 21, 85, 64, 45, 19])

```
In [49]: sns.distplot(final_array)
```

Out[49]:



```
In [25]: # Name of the Experiment : Missing and inappropriate data
# EX NO : 03 #Register Number : 230701059 #
Name : M N CHANDNI # Date : 19/08/2024
```

```
In [42]: import numpy as np
import pandas as pd
# Upload Hotel.csv and convert it into DataFrame
df = pd.read_csv("Hotel_Dataset.csv")
print("Original DataFrame:")
print(df)
# From the dataframe, identify duplicate rows (i.e., row 9)
print("\nChecking for duplicates:")
print(df.duplicated())

# The info() method prints information about the DataFrame, including the number of columns, column data types,
print("\nDataFrame Information:")
df.info()

# Remove duplicate rows
df.drop_duplicates(inplace=True)
print("\nDataFrame after removing duplicates:")
print(df)
# Reset index after dropping duplicate rows
print("\nResetting index after removing duplicates:")
df.reset_index(drop=True, inplace=True)
print(df)
# Use axis=1 to drop 'Age_Group.1' column from the DataFrame (if it exists)
df.drop(['Age_Group.1'], axis=1, inplace=True, errors='ignore')
print("\nDataFrame after dropping 'Age_Group.1' column:")
```



```

print(df)

# Correcting negative values in CustomerID, Bill, and EstimatedSalary using loc to avoid chained assignment
df.loc[df.CustomerID < 0, 'CustomerID'] = np.nan
df.loc[df.Bill < 0, 'Bill'] = np.nan
df.loc[df.EstimatedSalary < 0, 'EstimatedSalary'] = np.nan
print("\nDataFrame after replacing negative values with NaN:")
print(df)

# Replacing invalid 'NoOfPax' values (<1 or >20) with NaN using loc
df.loc[(df['NoOfPax'] < 1) | (df['NoOfPax'] > 20), 'NoOfPax'] = np.nan
print("\nDataFrame after replacing invalid 'NoOfPax' values with NaN:")
print(df)

# Show unique values of 'Age_Group', 'Hotel' and 'FoodPreference'
print("\nUnique values in 'Age_Group' column:")
print(df.Age_Group.unique())

print("\nUnique values in 'Hotel' column:")
print(df.Hotel.unique())
print("\nUnique values in 'FoodPreference' column:")
print(df.FoodPreference.unique())

# Replace incorrect or inconsistent values in 'Hotel' column using loc
df.loc[df.Hotel == 'Ibys', 'Hotel'] = 'Ibis'
print("\nDataFrame after replacing 'Ibys' with 'Ibis' in 'Hotel' column:")
print(df)

# Replace values in 'FoodPreference' column using loc
df.loc[df.FoodPreference.isin(['Vegetarian', 'veg']), 'FoodPreference'] = 'Veg'
df.loc[df.FoodPreference == 'non-Veg', 'FoodPreference'] = 'Non-Veg'
print("\nDataFrame after replacing inconsistent values in 'FoodPreference' column:")
print(df)

# Fill missing values in numerical columns with mean (for continuous) and median (for discrete) using loc
df.loc[:, 'EstimatedSalary'] = df['EstimatedSalary'].fillna(round(df['EstimatedSalary'].mean()))
df.loc[:, 'NoOfPax'] = df['NoOfPax'].fillna(round(df['NoOfPax'].median()))
df.loc[:, 'Rating(1-5)'] = df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()))
df.loc[:, 'Bill'] = df['Bill'].fillna(round(df['Bill'].mean()))

# Fill missing values in categorical columns (if needed) with the mode
df.loc[:, 'Age_Group'] = df['Age_Group'].fillna(df['Age_Group'].mode()[0])
df.loc[:, 'Hotel'] = df['Hotel'].fillna(df['Hotel'].mode()[0])
df.loc[:, 'FoodPreference'] = df['FoodPreference'].fillna(df['FoodPreference'].mode()[0])

# Display final cleaned DataFrame
print("\nFinal cleaned DataFrame:")
print(df)

# Save the cleaned DataFrame to a new CSV file
df.to_csv("Cleaned_Hotel_Dataset.csv", index=False)

```

OUTPUT

```

Original DataFrame:
  CustomerID Age_Group Rating(1-5) Hotel FoodPreference Bill \
0         1   20-25         4     Ibis      veg    1300
1         2   30-35         5  LemonTree  Non-Veg    2000
2         3   25-30         6    RedFox      Veg    1322
3         4   20-25        -1  LemonTree      Veg    1234
4         5    35+         3     Ibis  Vegetarian    989
5         6    35+         3     Ibys  Non-Veg    1909
6         7    35+         4    RedFox  Vegetarian    1000
7         8   20-25         7  LemonTree      Veg    2999
8         9   25-30         2     Ibis  Non-Veg    3456
9         9   25-30         2     Ibis  Non-Veg    3456
10        10   30-35         5    RedFox  non-Veg   -6755

  NoOfPax EstimatedSalary Age_Group.1
0         2         40000   20-25
1         3         59000   30-35
2         2         30000   25-30
3         2        120000   20-25
4         2         45000    35+
5         2        122220    35+
6        -1         21122    35+
7        -10        345673   20-25
8         3        -99999   25-30
9         3        -99999   25-30
10        4         87777   30-35

Checking for duplicates:
0  False
1  False
2  False

```

```

3      False 4      False 5      False 6      False 7      False 8      False 9      True 10     False dtype: bool DataFrame Information:
<class 'pandas.core.frame.DataFrame'> RangeIndex: 11 entries, 0 to 10 Data columns (total 9 columns):
# Column Non-Null Count Dtype ---
1 Age_Group 11 non-null object 2 Rating(1-5) 11 non-null int64 3 Hotel 11 non-null object
4 FoodPreference 11 non-null object 5 Bill 11 non-null int64 6 NoOfPax 11 non-null int64
7 EstimatedSalary 11 non-null int64 8 Age_Group.1 11 non-null object dtypes: int64(5), object(4)
memory usage: 924.0+ bytes DataFrame after removing duplicates: CustomerID Age_Group Rating(1-5)
5) Hotel FoodPreference Bill \ 0 1 20-25 4 Ibis veg 1300 1 2 30-35
35 5 LemonTree Non-Veg 2000 2 3 25-30 6 RedFox Veg 1322 3 4 20-25
25 -1 LemonTree Veg 1234 4 5 35+ 3 Ibis Vegetarian 989
5 6 35+ 3 Ibys Non-Veg 1909 6 7 35+ 4 RedFox Vegetarian 1000
7 8 20-25 7 LemonTree Veg 2999 8 9 25-30 2 Ibis Non-Veg 3456
10 10 30-35 5 RedFox non-Veg -6755 NoOfPax EstimatedSalary Age_Group.1
0 2 40000 20-25 1 3 59000 30-35 2 2 30000 25-30 3 2 120000 20-25
25 4 2 45000 35+ 5 2 122220 35+ 6 -1 21122 35+ 7 -10 345673 20-25
25 8 3 -99999 25-30 10 4 87777 30-35 Resetting index after removing duplicates:
CustomerID Age_Group Rating(1-5) Hotel FoodPreference Bill NoOfPax \ 0 1 20-25
25 4 Ibis veg 1300 2 1 2 30-35 5 LemonTree Non-Veg 2000 3 2 3 25-30
30 6 RedFox Veg 1322 2 3 4 20-25 -1 LemonTree Veg 1234 2
4 5 35+ 3 Ibis Vegetarian 989 2 5 6 35+ 3 Ibys Non-Veg 1909 2
6 7 35+ 4 RedFox Vegetarian 1000 -1 7 8 20-25
25 7 LemonTree Veg 2999 -10 8 9 25-30 2 Ibis Non-Veg 3456 3 9 10 30-35
35 5 RedFox non-Veg -6755 4 EstimatedSalary Age_Group.1 0 40000 20-25
1 59000 30-35 2 30000 25-30 3 120000 20-25 4 45000 35+
5 122220 35+ 6 21122 35+ 7 345673 20-25 8 -99999 25-30 9 87777 30-35
DataFrame after dropping 'Age_Group.1' column: CustomerID Age_Group Rating(1-5)
5) Hotel FoodPreference Bill NoOfPax \ 0 1 20-25 4 Ibis veg 1300 2 1 2 30-35
35 5 LemonTree Non-Veg 2000 3 2 3 25-30 6 RedFox Veg 1322 2 3 4 20-25
25 -1 LemonTree Veg 1234 2 4 5 35+ 3 Ibis Vegetarian 989 2 5 6 35+
3 Ibys Non-Veg 1909 2 6 7 35+ 4 RedFox Vegetarian 1000 -1 7 8 20-25 7
LemonTree Veg 2999 -10 8 9 25-30 2 Ibis Non-Veg 3456 3 9 10 30-35 5
RedFox non-Veg -6755 4 EstimatedSalary 0 40000 1 59000 2 30000 3 120000 4
45000 5 122220 6 21122 7 345673 8 -99999 9 87777 DataFrame after replacing negative
values with NaN: CustomerID Age_Group Rating(1-5) Hotel FoodPreference Bill \ 0 1.0 20-25 4 Ibis
veg 1300.0 1 2.0 30-35 5 LemonTree Non-Veg 2000.0 2 3.0 25-30 6 RedFox Veg
1322.0 3 4.0 20-25 -1 LemonTree Veg 1234.0 4 5.0 35+ 3 Ibis Vegetarian 989.0 5
6.0 35+ 3 Ibys Non-Veg 1909.0 6 7.0 35+ 4 RedFox Vegetarian 1000.0 7 8.0 20-25
25 7 LemonTree Veg 2999.0 8 9.0 25-30 2 Ibis Non-Veg 3456.0 9 10.0 30-35
5 RedFox non-Veg NaN NoOfPax EstimatedSalary 0 2 40000.0 1 3 59000.0 2 2
30000.0 3 2 120000.0 4 2 45000.0 5 2 122220.0 6 -1 21122.0 7 -10 345673.0 8
3 NaN 9 4 87777.0 DataFrame after replacing invalid 'NoOfPax' values with NaN: CustomerID Age_Group
Rating(1-5) Hotel FoodPreference Bill \ 0 1.0 20-25 4 Ibis veg 1300.0 1 2.0 30-35 5
LemonTree Non-Veg 2000.0 2 3.0 25-30 6 RedFox Veg 1322.0 3 4.0 20-25 -1
LemonTree Veg 1234.0 4 5.0 35+ 3 Ibis Vegetarian 989.0 5 6.0 35+ 3 Ibys
Non-Veg 1909.0 6 7.0 35+ 4 RedFox Vegetarian 1000.0 7 8.0 20-25 7 LemonTree
Veg 2999.0 8 9.0 25-30 2 Ibis Non-Veg 3456.0 9 10.0 30-35 5 RedFox non-Veg
NaN NoOfPax EstimatedSalary 0 2.0 40000.0 1 3.0 59000.0 2 2.0 30000.0 3 2.0 120000.0
4 2.0 45000.0 5 2.0 122220.0 6 NaN 21122.0 7 NaN 345673.0 8 3.0 NaN 9 4.0
87777.0 Unique values in 'Age_Group' column: ['20-25' '30-35' '25-30' '35+'] Unique values in 'Hotel' column: ['Ibis'
'LemonTree' 'RedFox' 'Ibys'] Unique values in 'FoodPreference' column: ['veg' 'Non-Veg' 'Veg' 'Vegetarian' 'non-Veg']
DataFrame after replacing 'Ibys' with 'Ibis' in 'Hotel' column: CustomerID Age_Group Rating(1-5) Hotel FoodPreference
Bill \ 0 1.0 20-25 4 Ibis veg 1300.0 1 2.0 30-35 5 LemonTree Non-Veg 2000.0

```

2	3.0	25-30	6	RedFox	Veg	1322.0	3	4.0	20-
25			-1	LemonTree			Veg	1234.0	
4		5.0	35+		3	Ibis	Vegetarian	989.0	
5		6.0	35+		3	Ibis	Non-Veg	1909.0	
6	7.0	35+	4	RedFox	Vegetarian	1000.0	7	8.0	20-
25		7	LemonTree		Veg	2999.0	8	9.0	25-
30		2	Ibis		Non-Veg	3456.0	9	10.0	30-
35	5	RedFox	non-Veg	NaN					

NoOfPax	EstimatedSalary
---------	-----------------

0	2.0	40000.0
1	3.0	59000.0
2	2.0	30000.0
3	2.0	120000.0
4	2.0	45000.0
5	2.0	122220.0
6	NaN	21122.0
7	NaN	345673.0
8	3.0	NaN
9	4.0	87777.0

DataFrame after replacing inconsistent values in 'FoodPreference' column:

CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill \
0	1.0	20-25	4	Ibis	Veg 1300.0
1	2.0	30-35	5	LemonTree	Non-Veg 2000.0
2	3.0	25-30	6	RedFox	Veg 1322.0
3	4.0	20-25	-1	LemonTree	Veg 1234.0
4	5.0	35+	3	Ibis	Veg 989.0
5	6.0	35+	3	Ibis	Non-Veg 1909.0
6	7.0	35+	4	RedFox	Veg 1000.0
7	8.0	20-25	7	LemonTree	Veg 2999.0
8	9.0	25-30	2	Ibis	Non-Veg 3456.0
9	10.0	30-35	5	RedFox	Non-Veg NaN

NoOfPax	EstimatedSalary
---------	-----------------

0	2.0	40000.0
1	3.0	59000.0
2	2.0	30000.0
3	2.0	120000.0
4	2.0	45000.0
5	2.0	122220.0
6	NaN	21122.0
7	NaN	345673.0
8	3.0	NaN
9	4.0	87777.0

Final cleaned DataFrame:

CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill \
0	1.0	20-25	4	Ibis	Veg 1300.0
1	2.0	30-35	5	LemonTree	Non-Veg 2000.0
2	3.0	25-30	6	RedFox	Veg 1322.0
3	4.0	20-25	-1	LemonTree	Veg 1234.0
4	5.0	35+	3	Ibis	Veg 989.0
5	6.0	35+	3	Ibis	Non-Veg 1909.0
6	7.0	35+	4	RedFox	Veg 1000.0
7	8.0	20-25	7	LemonTree	Veg 2999.0
8	9.0	25-30	2	Ibis	Non-Veg 3456.0
9	10.0	30-35	5	RedFox	Non-Veg 1801.0

NoOfPax	EstimatedSalary
---------	-----------------

0	2.0	40000.0
1	3.0	59000.0
2	2.0	30000.0
3	2.0	120000.0
4	2.0	45000.0
5	2.0	122220.0
6	2.0	21122.0
7	2.0	345673.0
8	3.0	96755.0
9	4.0	87777.0

```
In [38]: # Name of the Experiment : Data Preprocessing
# EX NO : 04
# Register Number : 230701048
# Name : AWINTHIKA SANTHANA
# Date : 27/08/2024
```

```
In [56]: import numpy as np
import pandas as pd
# Create a sample dataset
data = {
    'Country': ['France', 'Spain', 'Germany', 'Spain', 'Germany', 'France', 'Spain', 'France', 'France', 'France'],
    'Age': [44, 27, 30, 38, 40, 35, 38, 48, 50, 37],
    'Salary': [72000, 48000, 54000, 61000, 63778, 58000, 52000, 79000, 83000, 67000],
    'Purchased': ['No', 'Yes', 'No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes'] } # Create DataFrame df = pd.DataFrame(data)
# Display the original dataset print("Original DataFrame:") print(df) # Handling missing values (if any)
df['Country'] = df['Country'].fillna(df['Country'].mode()[0]) # Fill missing 'Country' with mode
df['Age'] = df['Age'].fillna(df['Age'].median()) # Fill missing 'Age' with median
df['Salary'] = df['Salary'].fillna(round(df['Salary'].mean())) # Fill missing 'Salary' with mean # One-hot encode the 'Country' column
df_encoded = pd.concat([pd.get_dummies(df['Country']), df[['Age', 'Salary', 'Purchased']]], axis=1)
# Handle the downcasting warning for 'Purchased' column
# Option 1: Setting the option to suppress downcasting warning pd.set_option('future.no_silent_downcasting', True)
# Replace the 'Purchased' column ('Yes'/'No' to 1/0)
df_encoded['Purchased'] = df_encoded['Purchased'].replace(['No', 'Yes'], [0, 1]) # Display the processed DataFrame
print("\nProcessed DataFrame:") print(df_encoded) # Additional Operations (to showcase more code)
# Calculate summary statistics summary_stats = df_encoded.describe()
# Group by countries and calculate mean of 'Age' and 'Salary'
country_grouped = df_encoded.groupby(['France', 'Germany', 'Spain']).agg({'Age': 'mean', 'Salary': 'mean'})
# Handle missing values (replacing 'Purchased' with the mode)
df_encoded['Purchased'] = df_encoded['Purchased'].fillna(df_encoded['Purchased'].mode()[0])
# Display the summary and grouped data print("\nSummary Statistics:") print(summary_stats)
print("\nCountry Grouped by Average Age and Salary:") print(country_grouped)
# Resetting the option to avoid future warnings pd.reset_option('future.no_silent_downcasting')
```

output

Original DataFrame: Country Age Salary Purchased
0 France 44 72000 No
1 Spain 27 48000 Yes
2 Germany 30 54000 No
3 Spain 38 61000 No
4 Germany 40 63778 Yes
5 France 35 58000 Yes
6 Spain 38 52000 No
7 France 48 79000 Yes
8 France 50 83000 No
9 France 37 67000 Yes

Processed DataFrame:

France Germany Spain Age Salary Purchased

0 True False False 44 72000 0

1 False False True 27 48000 1

2 False True False 30 54000 0

3 False False True 38 61000 0

4 False True False 40 63778 1

5 True False False 35 58000 1

6 False False True 38 52000 0

7 True False False 48 79000 1

8 True False False 50 83000 0

9 True False False 37 67000 1

Summary Statistics:

Age Salary

count 10.000000 10.000000

mean 38.700000 63777.800000

std 7.257946 11564.099406

min 27.000000 48000.000000

25% 35.500000 55000.000000

50% 38.000000 62389.000000

75% 43.000000 70750.000000

max 50.000000 83000.000000

Country Grouped by Average Age and Salary:

Age Salary

France Germany Spain

False False True 34.333333 53666.666667

True False 35.000000 58889.000000

True False False 42.800000 71800.000000

```
In [52]: # Name of the Experiment : EDA-Quantitative and Qualitative plots
# EX NO : 05
# Register Number : 230701048
# Name : awinthika santhanam
# Date : 03/09/2024
```

```
In [61]: import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
# Load the 'tips' dataset from seaborn
tips = sns.load_dataset('tips')
# Display the first few rows of the dataset
print(tips.head())
# Visualization 1: Displot with KDE for the 'total_bill' column
sns.displot(tips.total_bill, kde=True)
plt.title("Displot of Total Bill with KDE")
plt.show()
# Visualization 2: Displot without KDE for the 'total_bill' column
sns.displot(tips.total_bill, kde=False)
plt.title("Displot of Total Bill without KDE")
plt.show()
# Visualization 3: Jointplot for 'tip' vs 'total_bill'
sns.jointplot(x=tips.tip, y=tips.total_bill)
plt.title("Jointplot of Tip vs Total Bill")
plt.show()
# Visualization 4: Jointplot with regression line for 'tip' vs 'total_bill'
sns.jointplot(x=tips.tip, y=tips.total_bill, kind="reg")
plt.title("Jointplot with Regression of Tip vs Total Bill")
plt.show()
# Visualization 5: Jointplot with hexbin for 'tip' vs 'total_bill'
sns.jointplot(x=tips.tip, y=tips.total_bill, kind="hex")
```

```

plt.title("Hexbin Jointplot of Tip vs Total Bill")
plt.show()

# Visualization 6: Pairplot of all numerical columns
sns.pairplot(tips)
plt.title("Pairplot of Numerical Columns")
plt.show()

# Visualization 7: Pairplot with hue based on 'time'
sns.pairplot(tips, hue='time')
plt.title("Pairplot with Hue on Time")
plt.show()

# Visualization 8: Pairplot with hue based on 'day'
sns.pairplot(tips, hue='day')
plt.title("Pairplot with Hue on Day")
plt.show()

# Visualization 9: Heatmap of correlation matrix for numerical columns
sns.heatmap(tips.corr(numeric_only=True), annot=True)
plt.title("Heatmap of Correlation Matrix")
plt.show()

# Visualization 10: Boxplot for 'total_bill'
sns.boxplot(tips.total_bill)
plt.title("Boxplot of Total Bill")
plt.show()

# Visualization 11: Boxplot for 'tip'
sns.boxplot(tips.tip)
plt.title("Boxplot of Tip")
plt.show()

# Visualization 12: Countplot of 'day'
sns.countplot(tips.day)
plt.title("Countplot of Day")
plt.show()

# Visualization 13: Countplot of 'sex'
sns.countplot(tips.sex)
plt.title("Countplot of Sex")
plt.show()


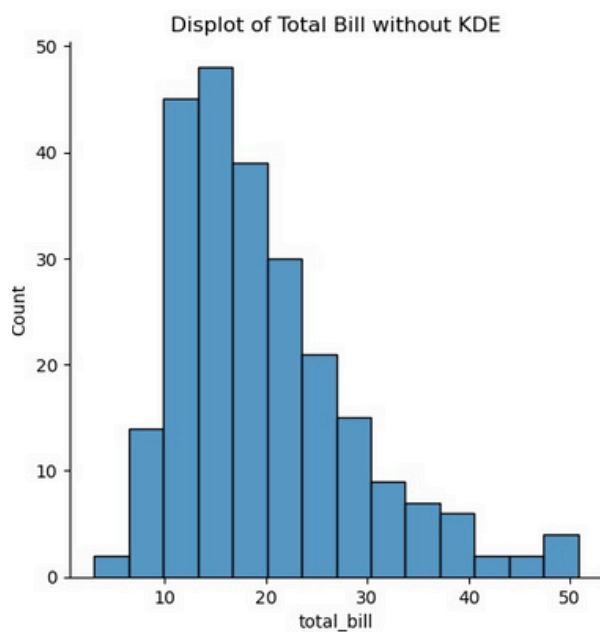

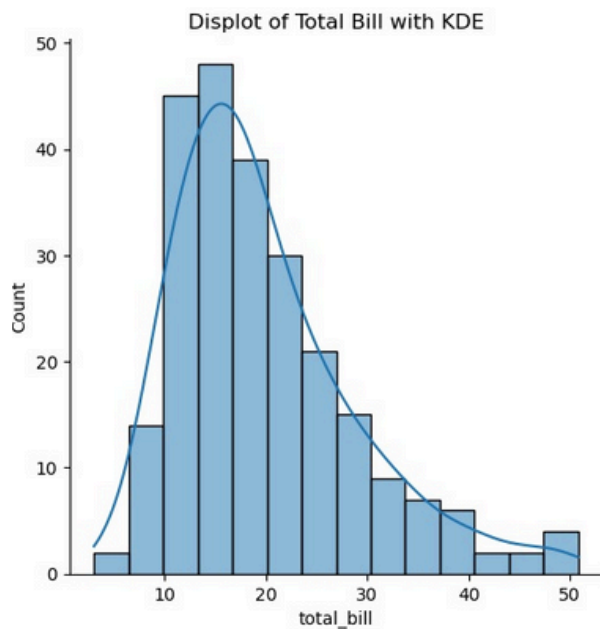
# Visualization 14: Pie chart of 'sex' value counts
tips.sex.value_counts().plot(kind='pie', autopct='%1.1f%%', startangle=90)
plt.title("Pie Chart of Sex Distribution")
plt.ylabel("") # Hide the 'sex' label
plt.show()

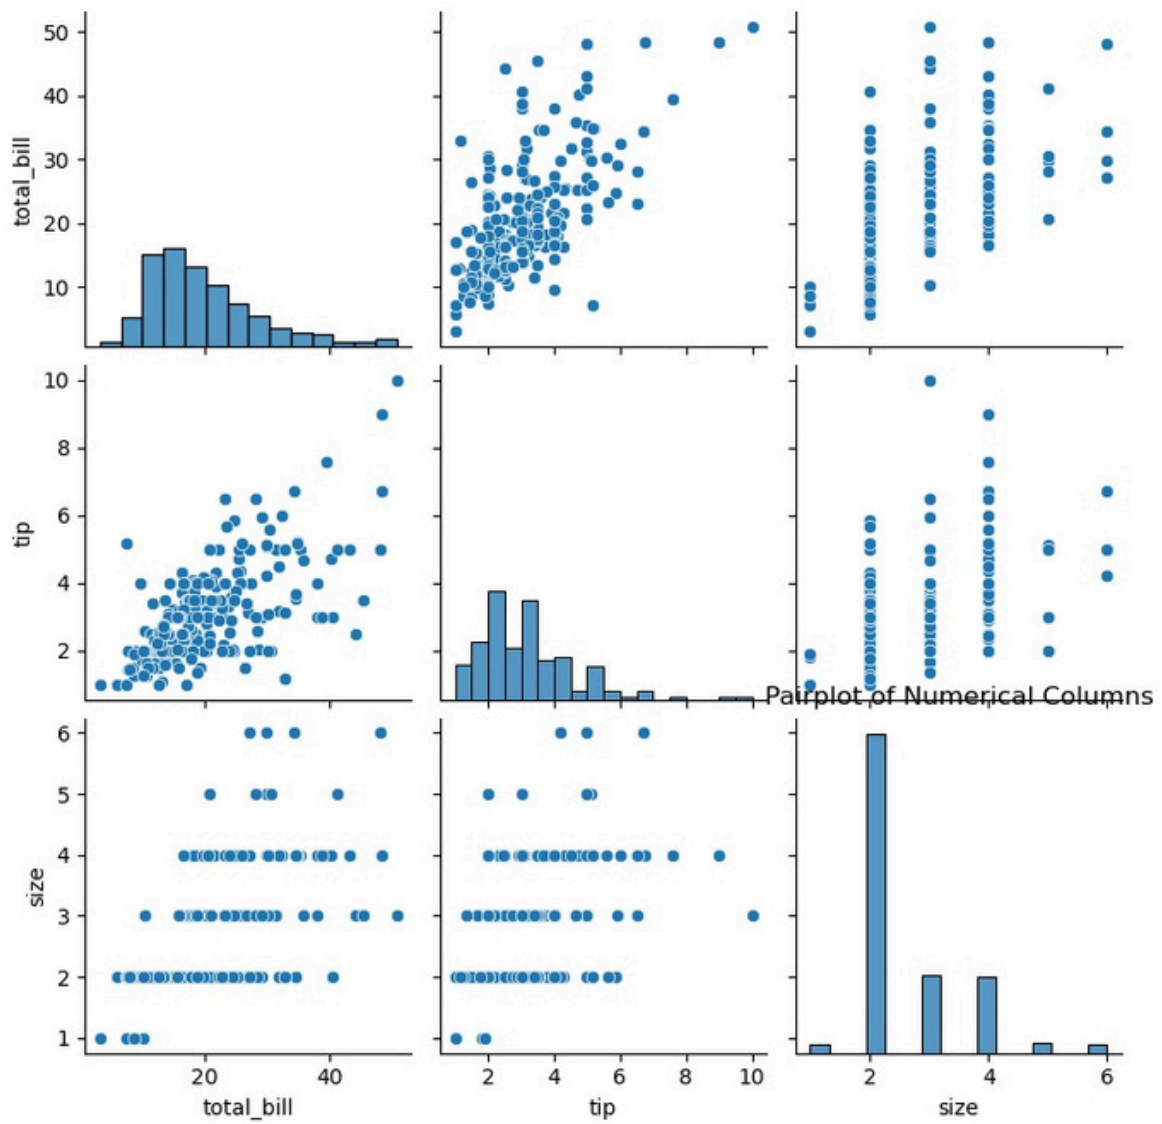

# Visualization 15: Bar chart of 'sex' value counts
tips.sex.value_counts().plot(kind='bar')
plt.title("Bar Chart of Sex Distribution")
plt.show()

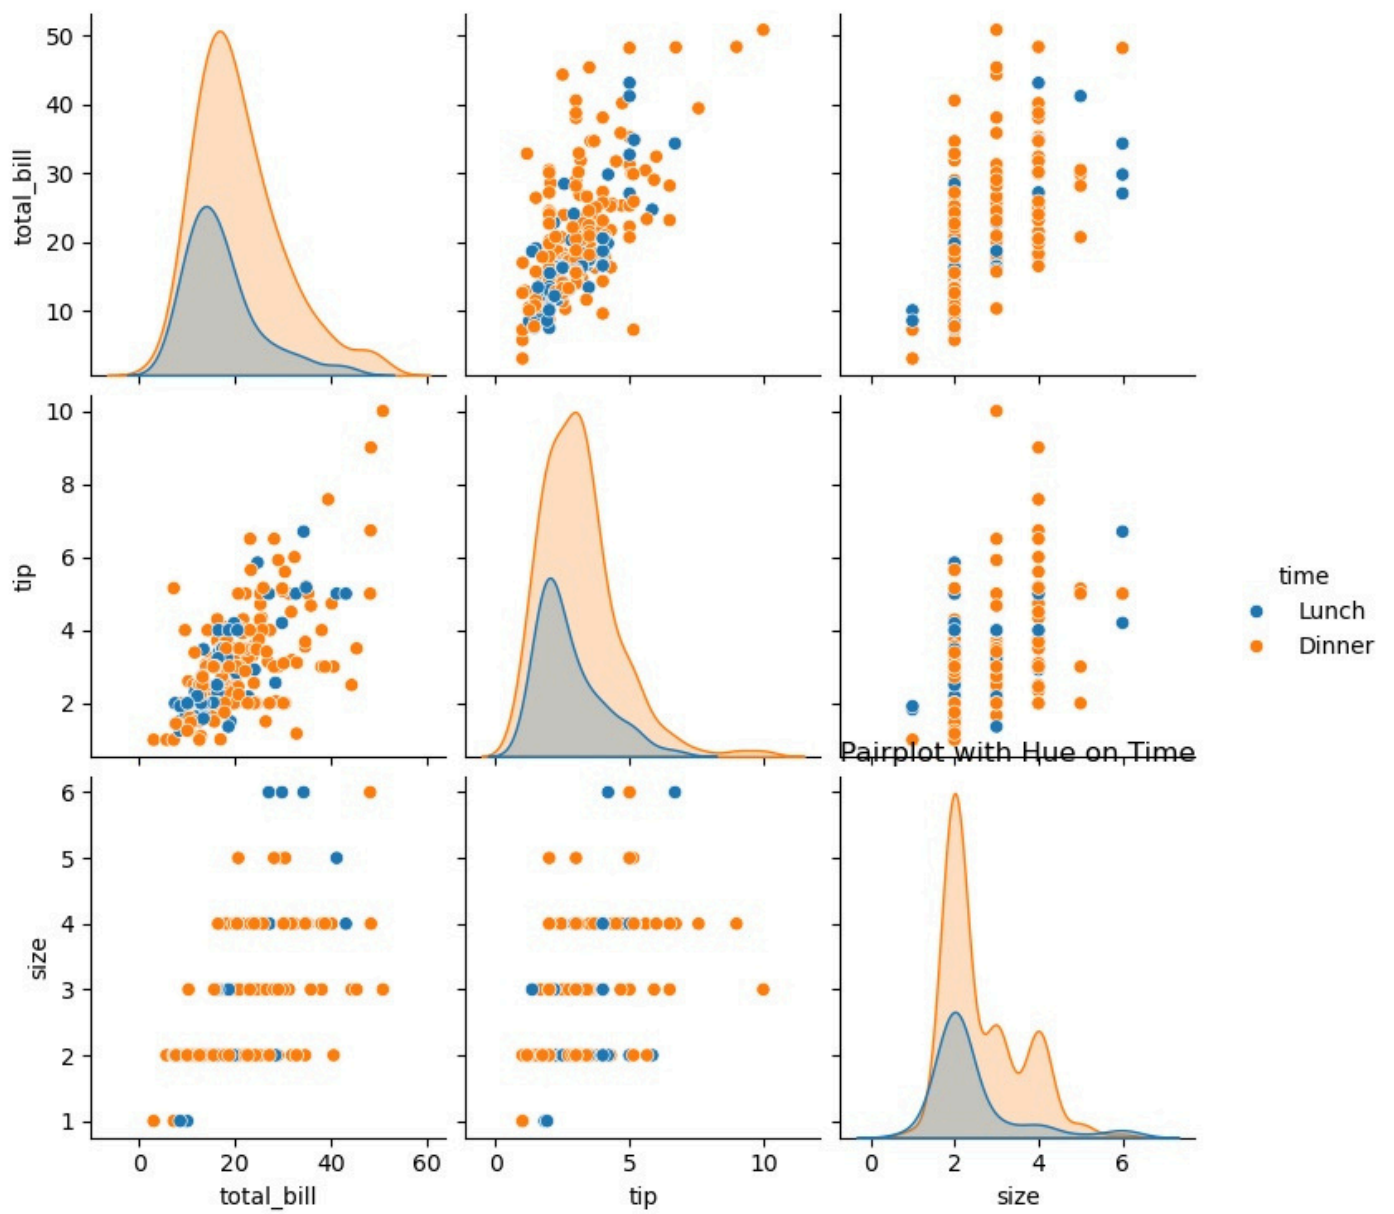
# Visualization 16: Countplot for 'day' based on 'time'=='Dinner'
sns.countplot(tips[tips.time=='Dinner']['day'])
plt.title("Countplot of Day for Dinner Time")
plt.show()

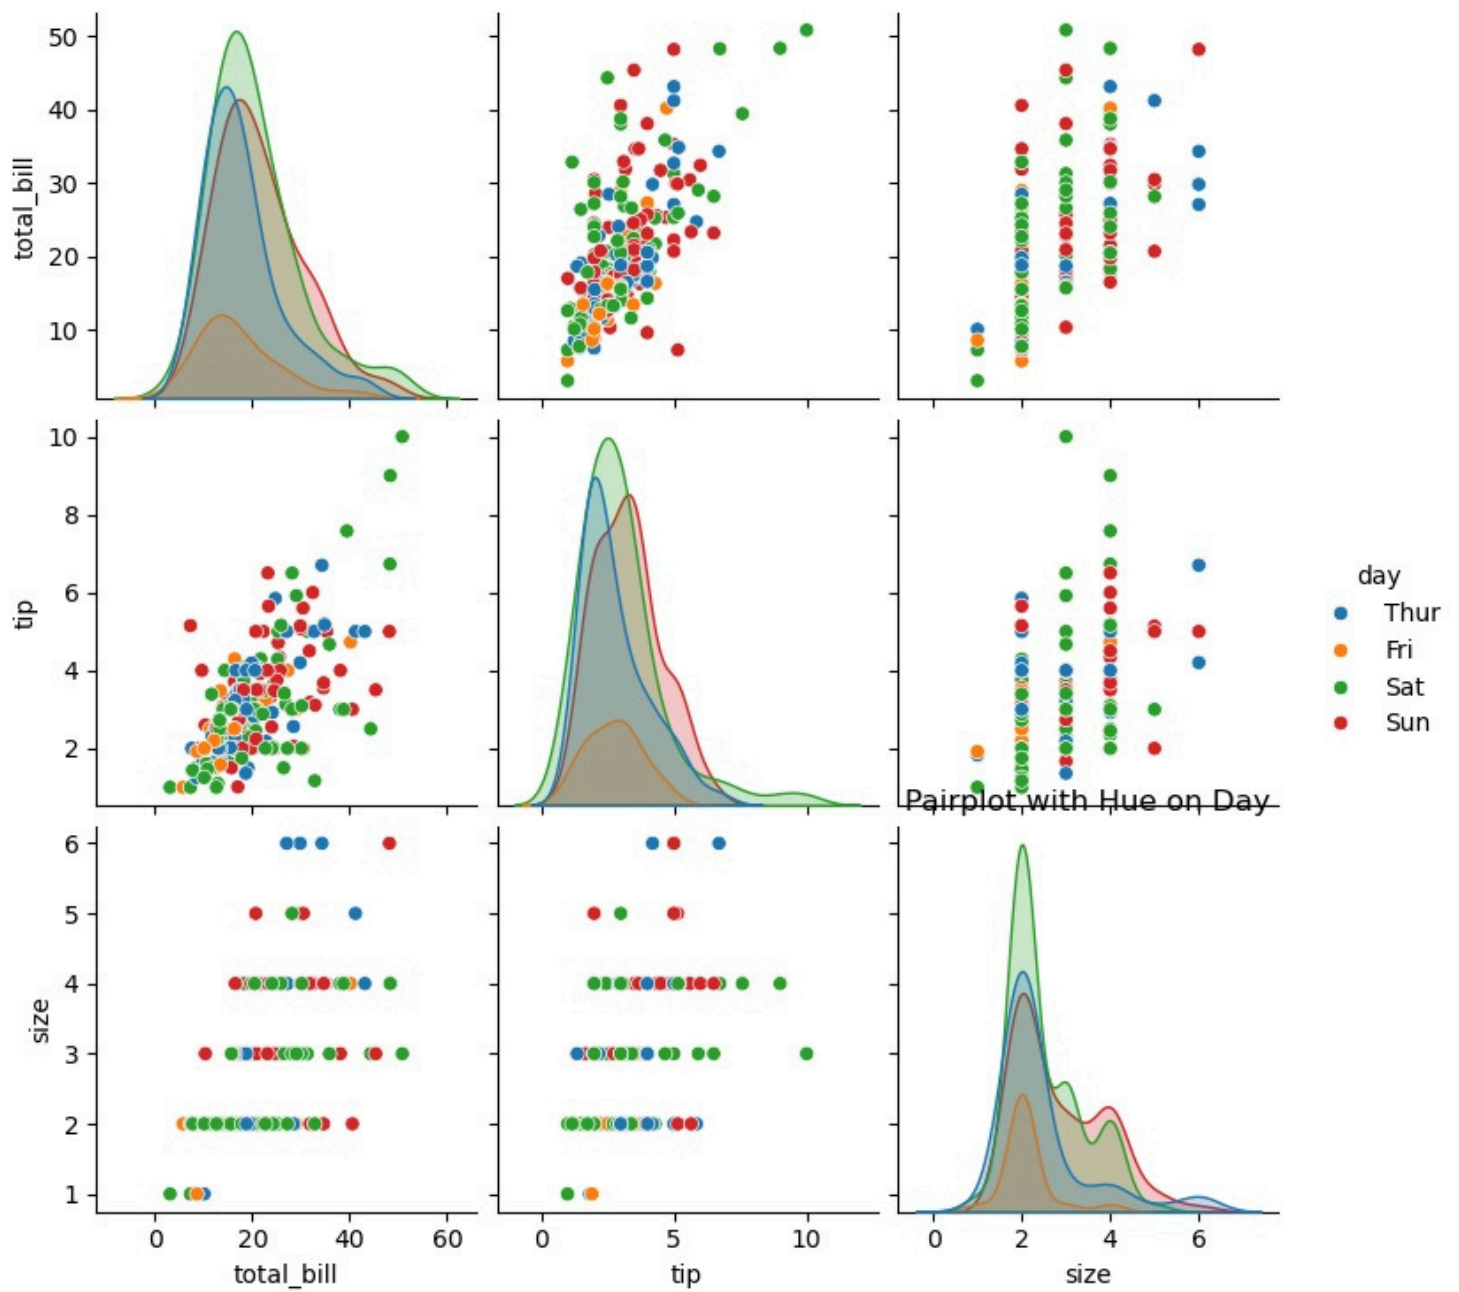
```



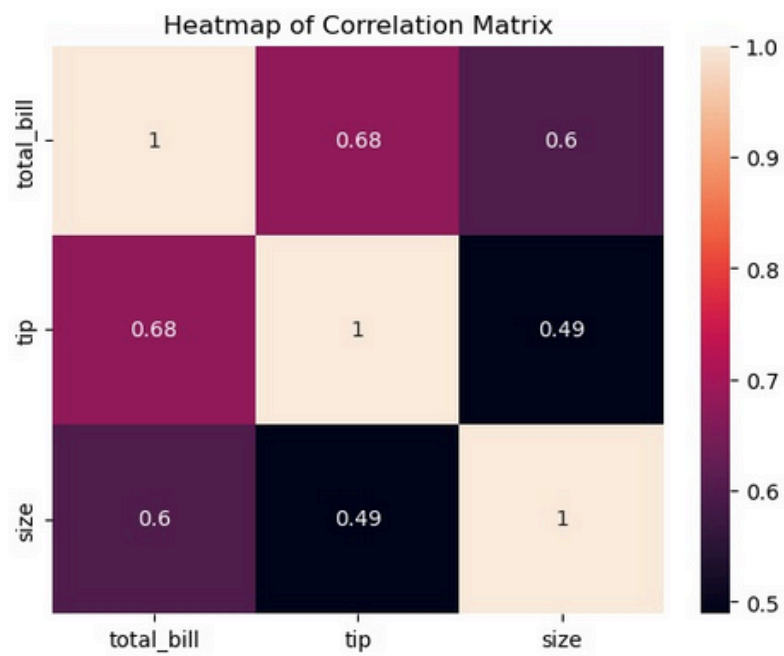
	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4


















Pie Chart of Sex Distribution



Bar Chart of Sex Distribution



Countplot of Day for Dinner Time






```
In [54]: # Name of the Experiment : Random Sampling and Sampling Distribution # EX NO : 06
#Register Number : 230701048 # Name : awinthika santhanam # Date :10/09/2024
```

```
In [63]: import numpy as np
import matplotlib.pyplot as plt
# Step 1: Generate a population (e.g., normal distribution)
population_mean = 50
population_std = 10
population_size = 100000
population = np.random.normal(population_mean, population_std, population_size)

# Step 2: Random sampling
sample_sizes = [30, 50, 100] # Different sample sizes to consider
num_samples = 1000 # Number of samples for each sample size
sample_means = {}
# Loop through each sample size
for size in sample_sizes:
    sample_means[size] = []
    for _ in range(num_samples):
        sample = np.random.choice(population, size=size, replace=False)
        sample_means[size].append(np.mean(sample))

# Step 3: Plotting sampling distributions
plt.figure(figsize=(12, 8))
# Loop through sample sizes and plot each distribution
for i, size in enumerate(sample_sizes):
    plt.subplot(len(sample_sizes), 1, i + 1)
    plt.hist(sample_means[size], bins=30, alpha=0.7, label=f'Sample Size {size}')
    plt.axvline(np.mean(population), color='red', linestyle='dashed', linewidth=1.5,
                label='Population Mean')
    plt.title(f'Sampling Distribution (Sample Size {size})')
    plt.xlabel('Sample Mean')
    plt.ylabel('Frequency')
    plt.legend()

# Adjust layout for better readability and show the plot
plt.tight_layout()
plt.show()
```




```
In [65]: #Name of the Experiment : Z-Test
# EX NO : 07 #
#Register Number : 230701048
#Name : awinthika santhanam
```

```
In [67]: import numpy as np
import scipy.stats as stats
# Define the sample data (hypothetical weights in grams)
sample_data = np.array([152, 148, 151, 149, 147, 153, 150, 148, 152,
                        149, 151, 150, 149, 152, 151, 148, 150, 152, 149,
                        150, 148, 153, 151, 150, 149, 152, 148, 151, 150, 153])
# Population mean under the null hypothesis
population_mean = 150
# Calculate sample statistics
sample_mean = np.mean(sample_data)
sample_std = np.std(sample_data, ddof=1) # Using sample standard deviation

# Number of observations
n = len(sample_data)
# Calculate the Z-statistic
z_statistic = (sample_mean - population_mean) / (sample_std / np.sqrt(n))
# Calculate the p-value (two-tailed test)
p_value = 2 * (1 - stats.norm.cdf(np.abs(z_statistic)))
# Print results
print(f"Sample Mean: {sample_mean:.2f}")
print(f"Z-Statistic: {z_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")
# Decision based on the significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: The average weight is significantly different from 150 grams.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in average weight from 150 gra
```

output

Sample Mean: 150.20
Z-Statistic: 0.6406
P-Value: 0.5218
Fail to reject the null hypothesis: There is no significant difference in average weight from 150 grams.

```
In [13]: # Name of the Experiment : T-Test
# EX NO : 08 #
Register Number : 230701048 # Name :
AWINTHIKA SANTHANAM
```

```
In [69]: import numpy as np
import scipy.stats as stats
# Set a random seed for reproducibility
np.random.seed(42)
# Generate hypothetical sample data (IQ scores)
sample_size = 25
sample_data = np.random.normal(loc=102, scale=15, size=sample_size) # Mean IQ of 102, SD of 15

# Population mean under the null hypothesis
population_mean = 100
# Calculate sample statistics
sample_mean = np.mean(sample_data)
sample_std = np.std(sample_data, ddof=1) # Using sample standard deviation

# Number of observations
n = len(sample_data)
# Calculate the T-statistic and p-value using a one-sample t-test
t_statistic, p_value = stats.ttest_1samp(sample_data, population_mean)
# Print results
print(f"Sample Mean: {sample_mean:.2f}")
print(f"T-Statistic: {t_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")
# Decision based on the significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: The average IQ score is significantly different from 100.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in average IQ score from 100.")
```

Sample Mean: 99.55

T-Statistic: -0.1577

P-Value: 0.8760

Fail to reject the null hypothesis: There is no significant difference in average IQ score from 100.

```
In [15]: # Name of the Experiment : Annova TEST
# EX NO : 09
# Register Number : 230701048
# Student Name : AWI THIKA SANTHANAM
```

```
In [71]: import numpy as np
import scipy.stats as stats
from statsmodels.stats.multicomp import pairwise_tukeyhsd

# Set a random seed for reproducibility
np.random.seed(42)
# Generate hypothetical growth data for three treatments (A, B, C)
n_plants = 25
# Growth data (in cm) for Treatment A, B, and C
growth_A = np.random.normal(loc=10, scale=2, size=n_plants)
growth_B = np.random.normal(loc=12, scale=3, size=n_plants)
growth_C = np.random.normal(loc=15, scale=2.5, size=n_plants)
# Combine all data into one array
all_data = np.concatenate([growth_A, growth_B, growth_C])
# Treatment labels for each group
treatment_labels = ['A'] * n_plants + ['B'] * n_plants + ['C'] * n_plants
# Perform one-way ANOVA
f_statistic, p_value = stats.f_oneway(growth_A, growth_B, growth_C)
# Print results
print("Treatment A Mean Growth:", np.mean(growth_A))
print("Treatment B Mean Growth:", np.mean(growth_B))
print("Treatment C Mean Growth:", np.mean(growth_C))
print()
print(f"F-Statistic: {f_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")

# Decision based on the significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in mean growth rates among the three tr
else:
    print("Fail to reject the null hypothesis: There is no significant difference in mean growth rates among the

# Additional: Post-hoc analysis (Tukey's HSD) if ANOVA is significant
if p_value < alpha:
    tukey_results = pairwise_tukeyhsd(all_data, treatment_labels, alpha=0.05)
    print("\nTukey's HSD Post-hoc Test:")
    print(tukey_results)
```

OUTPUT

Treatment A Mean Growth: 9.672983882683818
Treatment B Mean Growth: 11.137680744437432
Treatment C Mean Growth: 15.265234904828972

F-Statistic: 36.1214

P-Value: 0.0000

Reject the null hypothesis: There is a significant difference in mean growth rates among the three treatments.

Tukey's HSD Post-hoc Test:

Multiple Comparison of Means - Tukey HSD, FWER=0.05

=====

group1	group2	meandiff	p-adj	lower	upper	reject
--------	--------	----------	-------	-------	-------	--------

A	B	1.4647	0.0877	-0.1683	3.0977	False
A	C	5.5923	0.0	3.9593	7.2252	True
B	C	4.1276	0.0	2.4946	5.7605	True

```
In [17]: # Name of the Experiment : Feature Scaling
# EX NO : 10
# Register Number : 230701048
# Name : AWINTHIKA SANTHANAM
```

```
In [77]: import numpy as np
import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler, MinMaxScaler
from statsmodels.stats.multicomp import pairwise_tukeyhsd
import matplotlib.pyplot as plt
# Sample dataset
data = {
    'Country': ['France', 'Spain', 'Germany', 'Spain', 'Germany', 'France', 'Spain', 'France', 'France', 'France']
    'Age': [44, 27, 30, 38, 40, 35, 38, 48, 50, 37],
    'Salary': [72000, 48000, 54000, 61000, 65000, 58000, 52000, 79000, 83000, 67000],
    'Purchased': ['No', 'Yes', 'No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes']
}

# Create DataFrame
df = pd.DataFrame(data)
# Display the first few rows of the dataset
print("Original Data:")
print(df)

# Handle missing values
# Fill missing 'Country' with the mode (most frequent value)
df['Country'] = df['Country'].fillna(df['Country'].mode()[0])

# Separate features and labels
features = df.iloc[:, :-1].values
labels = df.iloc[:, -1].values

# Use SimpleImputer to handle missing values for 'Age' and 'Salary'
age_imputer = SimpleImputer(strategy="mean")
salary_imputer = SimpleImputer(strategy="mean")

# Impute missing values
features[:, 1] = age_imputer.fit_transform(features[:, [1]]).flatten()
features[:, 2] = salary_imputer.fit_transform(features[:, [2]]).flatten()

# OneHotEncoder for 'Country' column
oh = OneHotEncoder(sparse_output=False)
country_encoded = oh.fit_transform(features[:, [0]])

# Combine the encoded 'Country' values with the rest of the features
final_features = np.concatenate((country_encoded, features[:, 1:]), axis=1)
# Standardize the features using StandardScaler
scaler = StandardScaler()
standardized_features = scaler.fit_transform(final_features)

# Normalize the features using MinMaxScaler mms = MinMaxScaler(feature_range=(0, 1))
normalized_features = mms.fit_transform(standardized_features)

# Display the final processed features
print("\nProcessed Features (Standardized):")
print(standardized_features)

print("\nProcessed Features (Normalized):")
print(normalized_features)
# Plotting the processed data (just an example with a histogram for 'Salary')
plt.hist(df['Salary'], bins=10, color='skyblue', edgecolor='black')
plt.title('Salary Distribution')
plt.xlabel('Salary')
plt.ylabel('Frequency')
plt.show()

# Perform One-Way ANOVA to compare the mean 'Salary' across countries
from scipy import stats
f_stat, p_value = stats.f_oneway(df[df['Country'] == 'France']['Salary'],
                                df[df['Country'] == 'Spain']['Salary'],
                                df[df['Country'] == 'Germany']['Salary'])

print("\nANOVA Results:")
print(f"F-Statistic: {f_stat:.4f}, P-Value: {p_value:.4f}")
# Decision based on significance level
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in mean Salary across countries.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in mean Salary across countries")
# Perform Tukey's HSD test if ANOVA is significant
if p_value < alpha:
    tukey_results = pairwise_tukeyhsd(df['Salary'], df['Country'], alpha=0.05)
    print("\nTukey's HSD Post-hoc Test Results:")
    print(tukey_results)
```

OUTPUT

Original Data:


	Country	Age	Salary	Purchased
0	France	44	72000	No
1	Spain	27	48000	Yes
2	Germany	30	54000	No
3	Spain	38	61000	No
4	Germany	40	65000	Yes
5	France	35	58000	Yes
6	Spain	38	52000	No
7	France	48	79000	Yes
8	France	50	83000	No
9	France	37	67000	Yes

Processed Features (Standardized):

[1.	-0.5	-0.65465367	0.76973439	0.7379204]
[-1.	-0.5	1.52752523	-1.69922498	-1.44851041]
[-1.	2.	-0.65465367	-1.26352627	-0.90190271]
[-1.	-0.5	1.52752523	-0.10166303	-0.26419372]
[-1.	2.	-0.65465367	0.18880278	0.10021141]
[1.	-0.5	-0.65465367	-0.53736175	-0.53749758]
[-1.	-0.5	1.52752523	-0.10166303	-1.08410528]
[1.	-0.5	-0.65465367	1.35066601	1.37562939]
[1.	-0.5	-0.65465367	1.64113182	1.74003452]
[1.	-0.5	-0.65465367	-0.24689594	0.28241398]

Processed Features (Normalized):

[1.	0.	0.	0.73913043	0.68571429]
[0.	0.	1.	0.	0.]
[0.	1.	0.	0.13043478	0.17142857]
[0.	0.	1.	0.47826087	0.37142857]
[0.	1.	0.	0.56521739	0.48571429]
[1.	0.	0.	0.34782609	0.28571429]
[0.	0.	1.	0.47826087	0.11428571]
[1.	0.	0.	0.91304348	0.88571429]
[1.	0.	0.	1.	1.]
[1.	0.	0.	0.43478261	0.54285714]



ANOVA Results:
F-Statistic: 4.3100, P-Value: 0.0602
Fail to reject the null hypothesis: There is no significant difference in mean Salary across countries.

```
In [19]: # Name of the Experiment : Linear Regression
# EX NO : 11
# Register Number : 230701048
# Name : AWINTHIKA SANTHANAM
```

```
In [81]: # Importing necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

# Sample dataset
data = {
    'Country': ['France', 'Spain', 'Germany', 'Spain', 'Germany', 'France', 'Spain', 'France', 'France', 'France'],
    'Age': [44, 27, 30, 38, 40, 35, 38, 48, 50, 37],
    'Salary': [72000, 48000, 54000, 61000, 65000, 58000, 52000, 79000, 83000, 67000],
    'Purchased': ['No', 'Yes', 'No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes']
}

# Create DataFrame
df = pd.DataFrame(data)
# Handle missing values (if any)
df['Country'] = df['Country'].fillna(df['Country'].mode()[0]) # Filling missing countries
df['Age'] = df['Age'].fillna(df['Age'].mean()) # Filling missing age
df['Salary'] = df['Salary'].fillna(df['Salary'].mean()) # Filling missing salary
# Encoding 'Country' (categorical to numerical using OneHotEncoding)
from sklearn.preprocessing import OneHotEncoder
# Encode the 'Country' feature
encoder = OneHotEncoder(sparse_output=False) # Updated parameter
country_encoded = encoder.fit_transform(df[['Country']])


# Combine the encoded 'Country' with 'Age' and 'Salary'
X = np.concatenate((country_encoded, df[['Age', 'Salary']].values), axis=1)
# Convert 'Purchased' to numeric (0 for 'No', 1 for 'Yes')
df['Purchased'] = df['Purchased'].map({'No': 0, 'Yes': 1})
y = df['Purchased'].values

# Split the dataset into training and testing sets (80% training, 20% testing) X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.2, random_state=42)
# Feature scaling (standardization)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
# Initialize and train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)
# Make predictions
y_pred = model.predict(X_test)
# Evaluate the model
mse = mean_squared_error(y_test, y_pred) # Mean Squared Error
rmse = np.sqrt(mse) # Root Mean Squared Error
r2 = r2_score(y_test, y_pred) # R-squared score

# Output evaluation metrics
print(f"Mean Squared Error: {mse:.4f}")
print(f"Root Mean Squared Error: {rmse:.4f}")
print(f"R-squared: {r2:.4f}")
# Visualize the comparison of predicted vs actual values
plt.figure(figsize=(10, 6))
plt.scatter(range(len(y_test)), y_test, color='blue', label='Actual')
plt.scatter(range(len(y_test)), y_pred, color='red', label='Predicted')
plt.title('Actual vs Predicted Purchase')
plt.xlabel('Index')
plt.ylabel('Purchased (0/1)')
plt.legend()
plt.show()

# Optional: Visualizing the regression line for Salary vs Purchased
plt.figure(figsize=(10, 6))
plt.scatter(df['Salary'], df['Purchased'], color='blue')
plt.plot(df['Salary'], model.predict(scaler.transform(np.concatenate((encoder.transform(df[['Country']]), df[['A
plt.title('Salary vs Purchased')
plt.xlabel('Salary')
plt.ylabel('Purchased')
plt.show()
```

Mean Squared Error: 1.0261
Root Mean Squared Error: 1.0130
R-squared: -3.1044






```
In [21]: # Name of the Experiment : Logistic Regression # EX NO : 12
#Register Number : 230701048 # Name : AWINTHIKA SANTHANAM
```

```
In [101]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report
# Corrected dataset with equal-length lists
data = {
    "User ID": [15624510, 15810944, 15668575, 15603246, 15804002, 15683016, 15707098, 15686536, 15621310, 156823
15746732, 15680352, 15820022, 15636760, 15717341, 15755018, 15691863, 15706071, 15654296, 157550
    "Gender": ['Male', 'Male', 'Female', 'Female', 'Male', 'Female', 'Male', 'Female', 'Male', 'Female',
'Male', 'Female', 'Female', 'Male', 'Female', 'Male', 'Female', 'Male', 'Female', 'Male'],
    "Age": [19, 35, 26, 27, 19, 30, 35, 38, 28, 25, 35, 31, 35, 32, 34, 36, 46, 51, 50, 36],
    "EstimatedSalary": [19000, 20000, 43000, 57000, 76000, 85000, 150000, 60000, 62000, 55000,
90000, 50000, 58000, 45000, 80000, 33000, 41000, 23000, 20000, 33000],
    "Purchased": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0]
}

# Convert the dictionary to a DataFrame
df = pd.DataFrame(data)
# Display the DataFrame to ensure it's correct
print("Dataset:\n", df.head())
# Preprocessing the data
df['Gender'] = df['Gender'].map({'Male': 0, 'Female': 1}) # Encoding 'Gender'
# Features and labels
features = df[['Gender', 'Age', 'EstimatedSalary']].values
labels = df['Purchased'].values

# Split the dataset into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(features, labels, test_size=0.2, random_state=0)
# Standardize the features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)

# Train the Logistic Regression model
model = LogisticRegression()
model.fit(x_train, y_train)

# Evaluate the model
train_score = model.score(x_train, y_train)
test_score = model.score(x_test, y_test)

print(f"\nTraining Accuracy: {train_score:.4f}")
print(f"Testing Accuracy: {test_score:.4f}")
# Classification report
y_pred = model.predict(x_test)
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Predicting on the entire dataset (for the sake of example)
y_pred_full = model.predict(features)
print("\nFull dataset predictions:\n", y_pred_full)
```

OUTPUT

Dataset:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

Training Accuracy: 0.9375

Testing Accuracy: 1.0000

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	3
1	1.00	1.00	1.00	1
accuracy			1.00	4
macro avg	1.00	1.00	1.00	4
weighted avg	1.00	1.00	1.00	4

Full dataset predictions:

[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]